

# ITCS\_3156\_Final\_Project\_AndrewGreen

May 6, 2025

## 1 ITCS-3156 Final Project Spring 2025

### 1.1 Name: Andrew Green

### 1.2 Project Name: Credit Card Fraud Detection using Machine Learning

```
[1]: # Import the necessary modules

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from collections import Counter
import itertools

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, precision_score, confusion_matrix, \
    recall_score, f1_score
```

### 1.3 Data Loading

```
[2]: # Load the csv file
# dataset are downloaded from : https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud

credit_df = pd.read_csv("creditcard.csv")
credit_df.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	
		V8	V9 ...	V21	V22	V23	V24	V25	\

```

0  0.098698  0.363787  ... -0.018307  0.277838 -0.110474  0.066928  0.128539
1  0.085102 -0.255425  ... -0.225775 -0.638672  0.101288 -0.339846  0.167170
2  0.247676 -1.514654  ...  0.247998  0.771679  0.909412 -0.689281 -0.327642
3  0.377436 -1.387024  ... -0.108300  0.005274 -0.190321 -1.175575  0.647376
4 -0.270533  0.817739  ... -0.009431  0.798278 -0.137458  0.141267 -0.206010

```

```

      V26      V27      V28  Amount  Class
0 -0.189115  0.133558 -0.021053   149.62      0
1  0.125895 -0.008983  0.014724     2.69      0
2 -0.139097 -0.055353 -0.059752   378.66      0
3 -0.221929  0.062723  0.061458   123.50      0
4  0.502292  0.219422  0.215153    69.99      0

```

[5 rows x 31 columns]

## 1.4 Exploratory Data Analysis

```
[3]: credit_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
---  -
0   Time    284807 non-null  float64
1   V1       284807 non-null  float64
2   V2       284807 non-null  float64
3   V3       284807 non-null  float64
4   V4       284807 non-null  float64
5   V5       284807 non-null  float64
6   V6       284807 non-null  float64
7   V7       284807 non-null  float64
8   V8       284807 non-null  float64
9   V9       284807 non-null  float64
10  V10      284807 non-null  float64
11  V11      284807 non-null  float64
12  V12      284807 non-null  float64
13  V13      284807 non-null  float64
14  V14      284807 non-null  float64
15  V15      284807 non-null  float64
16  V16      284807 non-null  float64
17  V17      284807 non-null  float64
18  V18      284807 non-null  float64
19  V19      284807 non-null  float64
20  V20      284807 non-null  float64
21  V21      284807 non-null  float64
22  V22      284807 non-null  float64
23  V23      284807 non-null  float64

```

```

24 V24      284807 non-null float64
25 V25      284807 non-null float64
26 V26      284807 non-null float64
27 V27      284807 non-null float64
28 V28      284807 non-null float64
29 Amount   284807 non-null float64
30 Class     284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```

[4]: # Check for null values

credit_df.isnull().values.any()

```

```

[4]: np.False_

```

```

[5]: credit_df["Amount"].describe()

```

```

[5]: count      284807.000000
     mean         88.349619
     std        250.120109
     min          0.000000
     25%          5.600000
     50%         22.000000
     75%        77.165000
     max       25691.160000
     Name: Amount, dtype: float64

```

```

[6]: non_fraud = len(credit_df[credit_df.Class == 0])
     fraud = len(credit_df[credit_df.Class == 1])
     fraud_percent = (fraud / (fraud + non_fraud)) * 100

     print("Number of Genuine transactions: ", non_fraud)
     print("Number of Fraud transactions: ", fraud)
     print("Percentage of Fraud transactions: {:.4f}".format(fraud_percent))

```

```

Number of Genuine transactions: 284315
Number of Fraud transactions: 492
Percentage of Fraud transactions: 0.1727

```

The Percentage of Fraud transactions is only 0.17%. We can see how imbalance the data is !

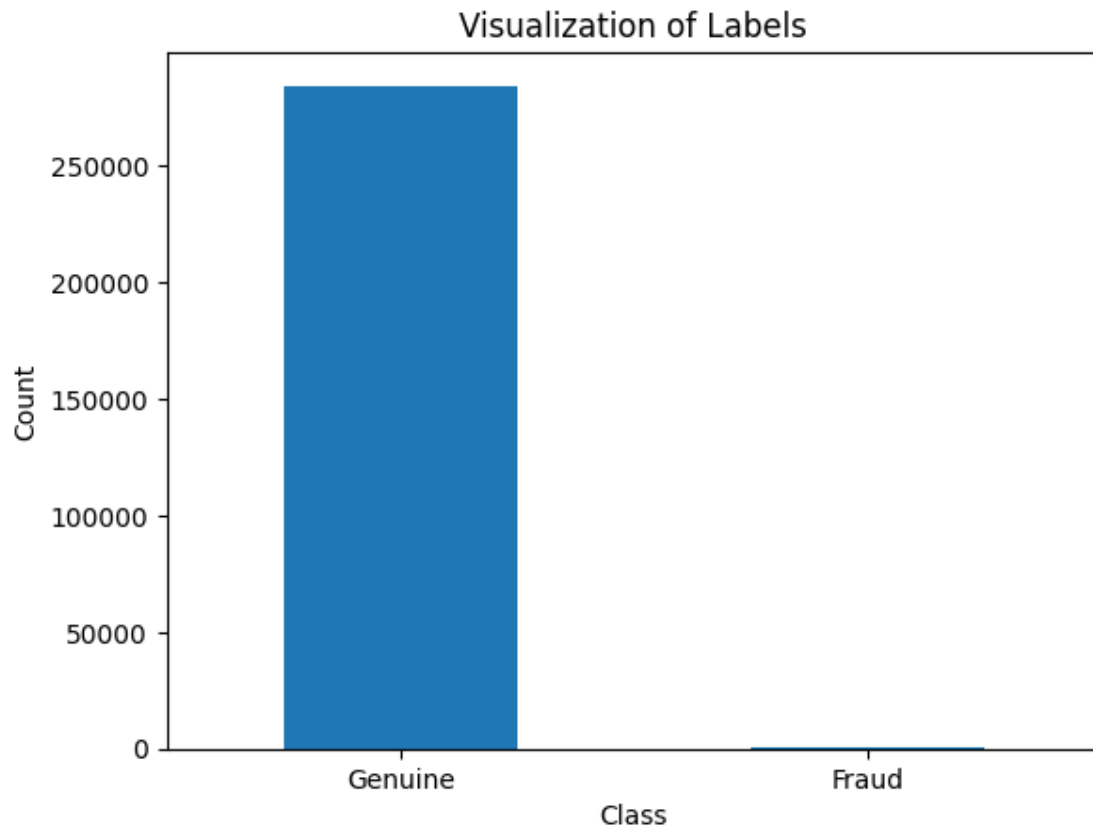
```

[7]: # Visualize the "Labels" column in our dataset

labels = ["Genuine", "Fraud"]
count_classes = credit_df.value_counts(credit_df['Class'], sort= True)
count_classes.plot(kind = "bar", rot = 0)
plt.title("Visualization of Labels")

```

```
plt.ylabel("Count")
plt.xticks(range(2), labels)
plt.show()
```



## 1.5 Data Preprocessing

```
[8]: # Perform Scaling on Amount
df2 = credit_df.copy()

scaler = StandardScaler()
df2["NormalizedAmount"] = scaler.fit_transform(df2["Amount"].values.reshape(-1, 1))
df2.drop(["Amount", "Time"], inplace=True, axis= 1)

y = df2["Class"]
X = df2.drop(["Class"], axis= 1)
```

```
[9]: y.head()
```

```
[9]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: Class, dtype: int64
```

```
[10]: # Split the data
      #(train_X, test_X, train_Y, test_Y) = train_test_split(X, Y, test_size= 0.3,
      ↪random_state= 42)
      X_trn, X_tst, y_trn, y_tst = train_test_split(X, y, train_size=.8,
      ↪random_state=42)
      X_trn, X_vld, y_trn, y_vld = train_test_split(X_trn, y_trn, train_size=.8,
      ↪random_state=42)

      print("Shape of X_trn: ", X_trn.shape)
      print("Shape of y_trn: ", y_trn.shape)
      print("Shape of X_vld: ", X_vld.shape)
      print("Shape of y_vld: ", y_vld.shape)
      print("Shape of X_tst: ", X_tst.shape)
      print("Shape of y_tst: ", y_tst.shape)
```

```
Shape of X_trn: (182276, 29)
Shape of y_trn: (182276,)
Shape of X_vld: (45569, 29)
Shape of y_vld: (45569,)
Shape of X_tst: (56962, 29)
Shape of y_tst: (56962,)
```

### 1.5.1 Helper function

```
[11]: # The below function is directly taken from the scikit-learn website to plot
      ↪the confusion matrix

      def plot_confusion_matrix(cm, classes, normalize=False, title='Confusion
      ↪Matrix', cmap=plt.cm.Blues):
          """
          This function prints and plots the confusion matrix.
          Normalization can be applied by setting `normalize=True`.
          """
          if normalize:
              cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
              print("Normalized confusion matrix")
          else:
              print('Confusion matrix, without normalization')
```

```

plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=0)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt), horizontalalignment="center",
    color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()

# The below function prints the following necessary metrics

def metrics(actuals, predictions):
    accuracy = accuracy_score(actuals, predictions)
    precision = precision_score(actuals, predictions)
    recall = recall_score(actuals, predictions)
    f1score = f1_score(actuals, predictions)
    print("Accuracy: {:.5f}".format(accuracy))
    print("Precision: {:.5f}".format(precision))
    print("Recall: {:.5f}".format(recall))
    print("F1-score: {:.5f}".format(f1score))

    return (accuracy, precision, recall, f1score)

```

## 1.6 Decision Tree Classifier

```

[12]: # Decision Tree Classifier
decision_tree = DecisionTreeClassifier()
decision_tree.fit(X_trn, y_trn)

predictions_dt = decision_tree.predict(X_tst)
decision_tree_score = decision_tree.score(X_tst, y_tst) * 100

```

```

[13]: # Plot confusion matrix for Decision Trees
confusion_matrix_dt = confusion_matrix(y_tst, predictions_dt.round())
print("Confusion Matrix - Decision Tree")
print(confusion_matrix_dt)

```

Confusion Matrix - Decision Tree


```

[[56838   26]
 [   22   76]]

```

```
[14]: plot_confusion_matrix(confusion_matrix_dt, classes=[0, 1], title= "Confusion_↵  
↵Matrix - Decision Tree")
```

Confusion matrix, without normalization



```
[15]: print("Evaluation of Decision Tree Model")  
print()  
metrics_dt = metrics(y_tst, predictions_dt.round())
```

Evaluation of Decision Tree Model

Accuracy: 0.99916  
Precision: 0.74510  
Recall: 0.77551  
F1-score: 0.76000

## 1.7 Neural Network Classifier

```
[16]: import keras
      from keras import layers
```

### 1.7.1 Neural Network Creation

```
[17]: model_kernel = keras.regularizers.l2(0.01)
      layer_unit = 128
      input_shape = (29,)

      model = keras.models.Sequential([
          keras.Input(shape=input_shape),
          layers.Flatten(),
          layers.Dense(layer_unit, activation='relu',
↪kernel_regularizer=model_kernel),
          layers.Dense(layer_unit, activation='relu',
↪kernel_regularizer=model_kernel),
          layers.Dropout(0.2),
          layers.Dense(1, activation='sigmoid')
      ])

      model.summary()
```

Model: "sequential"

Layer (type) ↪Param #	Output Shape	
flatten (Flatten) ↪ 0	(None, 29)	↪
dense (Dense) ↪3,840	(None, 128)	↪
dense_1 (Dense) ↪16,512	(None, 128)	↪
dropout (Dropout) ↪ 0	(None, 128)	↪
dense_2 (Dense) ↪129	(None, 1)	↪

Total params: 20,481 (80.00 KB)



Trainable params: 20,481 (80.00 KB)

Non-trainable params: 0 (0.00 B)

### 1.7.2 Model compile

```
[18]: model.compile(optimizer='adam',  
                  loss='binary_crossentropy',  
                  metrics=['accuracy'])
```

### 1.7.3 Train the model

```
[19]: # Train the model  
history = model.fit(X_trn, y_trn, epochs=20)
```

```
Epoch 1/20  
5697/5697          9s 1ms/step -  
accuracy: 0.9981 - loss: 0.1474  
Epoch 2/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9991 - loss: 0.0074  
Epoch 3/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9992 - loss: 0.0063  
Epoch 4/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9992 - loss: 0.0059  
Epoch 5/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9992 - loss: 0.0055  
Epoch 6/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9992 - loss: 0.0054  
Epoch 7/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9992 - loss: 0.0056  
Epoch 8/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9991 - loss: 0.0057  
Epoch 9/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9993 - loss: 0.0051  
Epoch 10/20  
5697/5697          8s 1ms/step -  
accuracy: 0.9992 - loss: 0.0050
```

```

Epoch 11/20
5697/5697          8s 1ms/step -
accuracy: 0.9991 - loss: 0.0056
Epoch 12/20
5697/5697          8s 1ms/step -
accuracy: 0.9992 - loss: 0.0052
Epoch 13/20
5697/5697          8s 1ms/step -
accuracy: 0.9991 - loss: 0.0054
Epoch 14/20
5697/5697          8s 1ms/step -
accuracy: 0.9993 - loss: 0.0048
Epoch 15/20
5697/5697          8s 1ms/step -
accuracy: 0.9992 - loss: 0.0053
Epoch 16/20
5697/5697          8s 1ms/step -
accuracy: 0.9993 - loss: 0.0045
Epoch 17/20
5697/5697          8s 1ms/step -
accuracy: 0.9992 - loss: 0.0050
Epoch 18/20
5697/5697          8s 1ms/step -
accuracy: 0.9993 - loss: 0.0043
Epoch 19/20
5697/5697          8s 1ms/step -
accuracy: 0.9992 - loss: 0.0054
Epoch 20/20
5697/5697          8s 1ms/step -
accuracy: 0.9993 - loss: 0.0047

```

```
[20]: model.predict(X_tst)
```

```
1781/1781          1s 627us/step
```

```
[20]: array([[7.6830745e-01],
           [2.7530853e-04],
           [3.8009480e-04],
           ...,
           [3.7469860e-04],
           [2.7530853e-04],
           [4.9429457e-04]], dtype=float32)
```

#### 1.7.4 Model Evaluate

```
[21]: loss, accuracy = model.evaluate(X_tst, y_tst, verbose=2)
print('Test accuracy:', accuracy)
```

```
1781/1781 - 2s - 1ms/step - accuracy: 0.9992 - loss: 0.0045
```

Test accuracy: 0.9991924166679382

## 1.8 Oversampling for imbalanced data

As mentioned above, we can clearly observe the imbalance on our credit dataset. When dataset is imbalanced, the model learns to favor the majority class, leading to a high accuracy score that can be misleading. This bias can result in the model failing to accurately identify or predict instances of the minority class, even if it's crucial for the overall purpose of the mode.

The consequences of a biased model can be significant, particularly in areas where accuracy and reliability are paramount. For example, in fraud detection, a model that favors detecting non-fraudulent transactions might miss actual fraudulent activities, leading to financial losses. In medical diagnosis, a biased model could misclassify a rare disease, potentially leading to delayed or incorrect treatment.

The dataset imbalance issue can be mitigated by various techniques, including Resampling, Class Weighting, Algorithm Selection and Adjusting Thresholding.

In this project, we will test the Resampling technique by oversampling the minority class (creating synthetic data) to balance the dataset.

```
[22]: #from imblearn.over_sampling import SMOTE
from imblearn.over_sampling import RandomOverSampler

ros = RandomOverSampler(sampling_strategy='auto',random_state=0)

X_resampled, y_resampled = ros.fit_resample(X, y)

print("Resampled shape of X: ", X_resampled.shape)
print("Resampled shape of y: ", y_resampled.shape)

value_counts = Counter(y_resampled)
print(value_counts)
```

Resampled shape of X: (568630, 29)

Resampled shape of y: (568630,)

Counter({0: 284315, 1: 284315})

### 1.8.1 Get train and test data from resampled data

```
[23]: # Split the data
X_trn_re, X_tst_re, y_trn_re, y_tst_re = train_test_split(X_resampled,
    ↪y_resampled, train_size=.8, random_state=42)
X_trn_re, X_vld_re, y_trn_re, y_vld_re = train_test_split(X_trn_re, y_trn_re,
    ↪train_size=.8, random_state=42)

print("Resampled shape of X_trn_re: ", X_trn_re.shape)
print("Resampled shape of y_trn_re: ", y_trn_re.shape)
print("Resampled shape of X_vld_re: ", X_vld_re.shape)
print("Resampled shape of y_vld_re: ", y_vld_re.shape)
```

```
print("Resampled shape of X_tst_re: ", X_vld_re.shape)
print("Resampled shape of y_tst_re: ", y_tst_re.shape)
```

```
Resampled shape of X_trn_re: (363923, 29)
Resampled shape of y_trn_re: (363923,)
Resampled shape of X_vld_re: (90981, 29)
Resampled shape of y_vld_re: (90981,)
Resampled shape of X_tst_re: (90981, 29)
Resampled shape of y_tst_re: (113726,)
```

## 1.9 Decision Tree Classifier Based on oversampled data

```
[24]: # Decision Tree Classifier
decision_tree.fit(X_trn_re, y_trn_re)

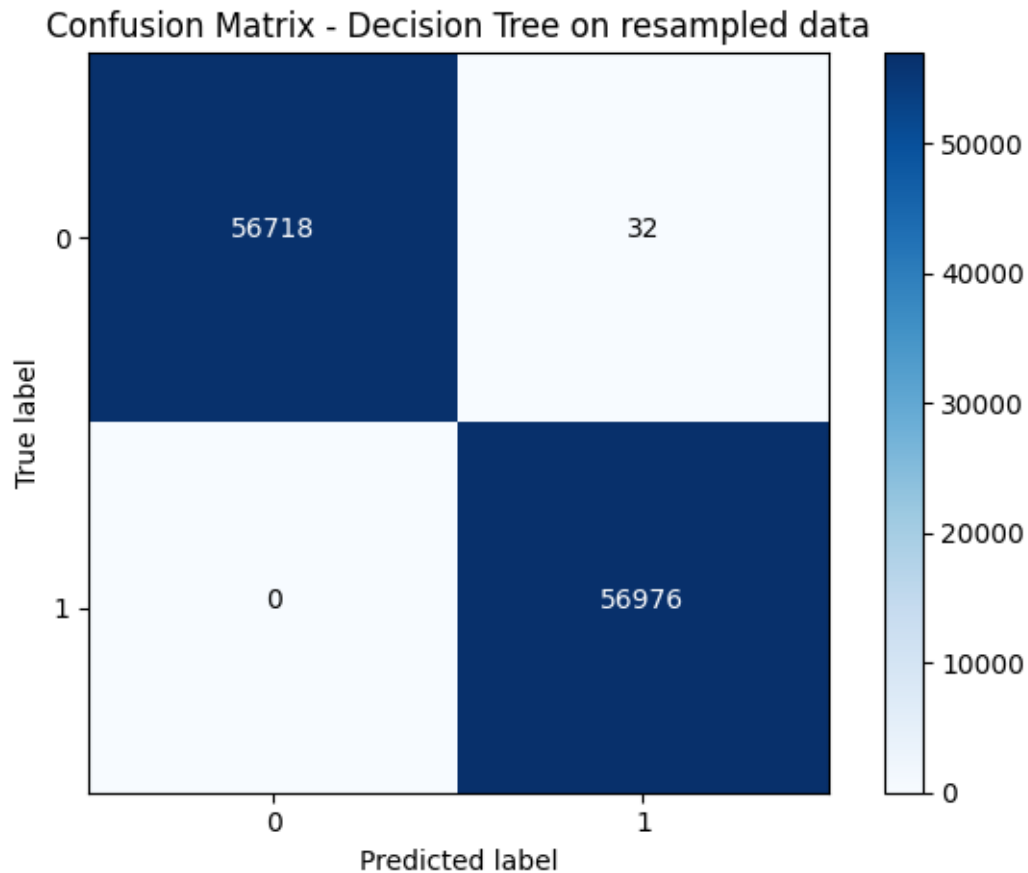
predictions_dt2 = decision_tree.predict(X_tst_re)
decision_tree_score2 = decision_tree.score(X_tst_re, y_tst_re) * 100
```

```
[25]: # Plot confusion matrix for Decision Trees
confusion_matrix_dt2 = confusion_matrix(y_tst_re, predictions_dt2.round())
print("Confusion Matrix - Decision Tree")
print(confusion_matrix_dt2)
```

```
Confusion Matrix - Decision Tree
[[56718   32]
 [    0 56976]]
```

```
[26]: plot_confusion_matrix(confusion_matrix_dt2, classes=[0, 1], title= "Confusion_
↪Matrix - Decision Tree on resampled data")
```

Confusion matrix, without normalization



```
[27]: print("Evaluation of Decision Tree Model over resampled data")
      print()
      metrics_dt2 = metrics(y_tst_re, predictions_dt2.round())
```

Evaluation of Decision Tree Model over resampled data

Accuracy: 0.99972  
Precision: 0.99944  
Recall: 1.00000  
F1-score: 0.99972

### 1.10 Neural Network Classifier over resampled data

```
[28]: model2 = keras.models.Sequential([
      keras.Input(shape=input_shape),
      layers.Flatten(),
      layers.Dense(layer_unit, activation='relu',
      ↪kernel_regularizer=model_kernel),
```

```

        layers.Dense(layer_unit, activation='relu',
↪kernel_regularizer=model_kernel),
        layers.Dropout(0.2),
        layers.Dense(1, activation='sigmoid')
])

model2.summary()

```

Model: "sequential\_1"

Layer (type) ↪Param #	Output Shape	
flatten_1 (Flatten) ↪ 0	(None, 29)	
dense_3 (Dense) ↪3,840	(None, 128)	
dense_4 (Dense) ↪16,512	(None, 128)	
dropout_1 (Dropout) ↪ 0	(None, 128)	
dense_5 (Dense) ↪129	(None, 1)	

Total params: 20,481 (80.00 KB)

Trainable params: 20,481 (80.00 KB)

Non-trainable params: 0 (0.00 B)

```

[29]: model2.compile(optimizer='adam',
                    loss='binary_crossentropy',
                    metrics=['accuracy'])

```

```

[30]: # Train the model
      history2 = model2.fit(X_trn_re, y_trn_re, epochs=20)

```

Epoch 1/20  
11373/11373                      17s 1ms/step

```

- accuracy: 0.9585 - loss: 0.2461
Epoch 2/20
11373/11373          16s 1ms/step
- accuracy: 0.9876 - loss: 0.0733
Epoch 3/20
11373/11373          16s 1ms/step
- accuracy: 0.9903 - loss: 0.0603
Epoch 4/20
11373/11373          16s 1ms/step
- accuracy: 0.9911 - loss: 0.0565
Epoch 5/20
11373/11373          16s 1ms/step
- accuracy: 0.9917 - loss: 0.0528
Epoch 6/20
11373/11373          16s 1ms/step
- accuracy: 0.9912 - loss: 0.0522
Epoch 7/20
11373/11373          17s 1ms/step
- accuracy: 0.9917 - loss: 0.0511
Epoch 8/20
11373/11373          17s 1ms/step
- accuracy: 0.9921 - loss: 0.0498
Epoch 9/20
11373/11373          16s 1ms/step
- accuracy: 0.9922 - loss: 0.0485
Epoch 10/20
11373/11373          17s 1ms/step
- accuracy: 0.9921 - loss: 0.0498
Epoch 11/20
11373/11373          17s 1ms/step
- accuracy: 0.9924 - loss: 0.0484
Epoch 12/20
11373/11373          17s 1ms/step
- accuracy: 0.9923 - loss: 0.0482
Epoch 13/20
11373/11373          17s 1ms/step
- accuracy: 0.9924 - loss: 0.0482
Epoch 14/20
11373/11373          17s 1ms/step
- accuracy: 0.9923 - loss: 0.0475
Epoch 15/20
11373/11373          16s 1ms/step
- accuracy: 0.9928 - loss: 0.0472
Epoch 16/20
11373/11373          17s 1ms/step
- accuracy: 0.9929 - loss: 0.0462
Epoch 17/20
11373/11373          17s 1ms/step

```

```
- accuracy: 0.9933 - loss: 0.0450
Epoch 18/20
11373/11373          17s 1ms/step
- accuracy: 0.9932 - loss: 0.0451
Epoch 19/20
11373/11373          17s 1ms/step
- accuracy: 0.9925 - loss: 0.0465
Epoch 20/20
11373/11373          17s 1ms/step
- accuracy: 0.9929 - loss: 0.0458
```

```
[31]: model2.predict(X_tst_re)
```

```
3554/3554          2s 620us/step
```

```
[31]: array([[9.9987364e-01],
            [9.9924558e-01],
            [6.4296324e-09],
            ...,
            [9.9922049e-01],
            [9.8870695e-01],
            [1.2334959e-07]], dtype=float32)
```

```
[32]: loss2, accuracy2 = model2.evaluate(X_tst_re, y_tst_re, verbose=2)
print('Test accuracy after resampling:', accuracy2)
```

```
3554/3554 - 4s - 1ms/step - accuracy: 0.9952 - loss: 0.0388
Test accuracy after resampling: 0.9952077865600586
```

## 1.11 Summary

Now it is evident that after addressing the class imbalance problem, our Decision Tree classifier performs far better than Decision Tree classifier without oversampling. At the same time, we see that the accuracy of Neural Network with oversampling slightly degenerate.

```
[ ]:
```